



Experiment No. 7
Apply Dimensionality Reduction on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:
Date of Submission:



Aim: Apply Dimensionality Reduction on Adult Census Income Dataset and analyze the performance of the model.

Objective: Able to perform various feature engineering tasks, perform dimensionality reduction on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

Theory:

In machine learning classification problems, there are often too many factors on the basis of which the final classification is done. These factors are basically variables called features. The higher the number of features, the harder it gets to visualize the training set and then work on it. Sometimes, most of these features are correlated, and hence redundant. This is where dimensionality reduction algorithms come into play. Dimensionality reduction is the process of reducing the number of random variables under consideration, by obtaining a set of principal variables. It can be divided into feature selection and feature extraction.

Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.



education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

Code:

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/"
# directory
# For example, running this (by clicking run or pressing Shift+Enter)
# will list all files under the input directory
```

```
import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
df=pd.read_csv("/content/adult.csv")
```

```
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status \
0	90	?	77053	HS-grad	9	Widowed
1	82	Private	132870	HS-grad	9	Widowed
2	66	?	186061	Some-college	10	Widowed
3	54	Private	140359	7th-8th	4	Divorced
4	41	Private	264663	Some-college	10	Separated

	capital.loss	hours.per.week	native.country	income
0	4356	40	United-States	<=50K
1	4356	18	United-States	<=50K
2	4356	40	United-States	<=50K
3	3900	40	United-States	<=50K
4	3900	40	United-States	<=50K

```
df.describe().T
```

	count	mean	std	min
25% \				
age	32561.0	38.581647	13.640433	17.0
28.0				
fnlwgt	32561.0	189778.366512	105549.977697	12285.0

```

117827.0
education.num    32561.0    10.080679    2.572720    1.0
9.0
capital.gain     32561.0    1077.648844    7385.292085    0.0
0.0
capital.loss     32561.0    87.303830    402.960219    0.0
0.0
hours.per.week   32561.0    40.437456    12.347429    1.0
40.0

```

```

          50%    75%    max
age          37.0    48.0    90.0
fnlwgt      178356.0  237051.0  1484705.0
education.num    10.0    12.0    16.0
capital.gain      0.0     0.0   99999.0
capital.loss      0.0     0.0   4356.0
hours.per.week    40.0    45.0    99.0

```

```
df.shape
```

```
(32561, 15)
```

```
df.columns
```

```

Index(['age', 'workclass', 'fnlwgt', 'education', 'education.num',
       'marital.status', 'occupation', 'relationship', 'race', 'sex',
       'capital.gain', 'capital.loss', 'hours.per.week',
       'native.country',
       'income'],
      dtype='object')

```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 32561 entries, 0 to 32560
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	age	32561 non-null	int64
1	workclass	32561 non-null	object
2	fnlwgt	32561 non-null	int64
3	education	32561 non-null	object
4	education.num	32561 non-null	int64
5	marital.status	32561 non-null	object
6	occupation	32561 non-null	object
7	relationship	32561 non-null	object
8	race	32561 non-null	object
9	sex	32561 non-null	object
10	capital.gain	32561 non-null	int64
11	capital.loss	32561 non-null	int64
12	hours.per.week	32561 non-null	int64

```
13 native.country 32561 non-null object
14 income          32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
df[df == '?'] = np.nan
```

```
df.isnull().sum()
```

```
for col in ['workclass', 'occupation', 'native.country']:
    df[col].fillna(df[col].mode()[0], inplace=True)

df.isnull().sum()
```

```
X = df.drop(['income'], axis=1)
```

```
y = df['income']
```

```

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 0)

from sklearn import preprocessing

categorical = ['workclass', 'education', 'marital.status',
'occupation', 'relationship', 'race', 'sex', 'native.country']
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X_train = pd.DataFrame(scaler.fit_transform(X_train), columns =
X.columns)

X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
X_train.head()

```

	age	workclass	fnlwgt	education	education.num	
marital.status \						
0	0.101484	2.600478	-1.494279	-0.332263	1.133894	-
0.402341						
1	0.028248	-1.884720	0.438778	0.184396	-0.423425	-
0.402341						
2	0.247956	-0.090641	0.045292	1.217715	-0.034095	
0.926666						
3	-0.850587	-1.884720	0.793152	0.184396	-0.423425	
0.926666						
4	-0.044989	-2.781760	-0.853275	0.442726	1.523223	-
0.402341						

	occupation	relationship	race	sex	capital.gain	
capital.loss \						
0	-0.782234	2.214196	0.39298	-1.430470	-0.145189	-
0.217407						
1	-0.026696	-0.899410	0.39298	0.699071	-0.145189	-
0.217407						
2	-0.782234	-0.276689	0.39298	-1.430470	-0.145189	-
0.217407						
3	-0.530388	0.968753	0.39298	0.699071	-0.145189	-
0.217407						
4	-0.782234	-0.899410	0.39298	0.699071	-0.145189	-
0.217407						

hours.per.week native.country

0	-1.662414	0.262317
1	-0.200753	0.262317
2	-0.038346	0.262317
3	-0.038346	0.262317
4	-0.038346	0.262317

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
LR = LogisticRegression()
LR.fit(X_train, y_train)
```

```
LogisticRegression()
```

```
y_pred = LR.predict(X_test)
```

```
accuracy_score(y_test, y_pred)
```

```
0.8216808271061521
```

```
from sklearn.decomposition import PCA
pca = PCA()
```

```
X_train = pca.fit_transform(X_train)
pca.explained_variance_ratio_
```

```
array([0.14757168, 0.10182915, 0.08147199, 0.07880174, 0.07463545,
        0.07274281, 0.07009602, 0.06750902, 0.0647268 , 0.06131155,
        0.06084207, 0.04839584, 0.04265038, 0.02741548])
```

```
X = df.drop(['income', 'native.country'], axis=1)
y = df['income']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 0)
```

```
categorical = ['workclass', 'education', 'marital.status',
'occupation', 'relationship', 'race', 'sex']
```

```
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])
```

```
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns =
X.columns)
```

```
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
```

```
LR1 = LogisticRegression()
LR1.fit(X_train, y_train)
```

```
LogisticRegression()
```



```

y_pred = LR1.predict(X_test)
accuracy_score(y_test, y_pred)

0.8212713686150066

X = df.drop(['income', 'native.country', 'hours.per.week'], axis=1)
y = df['income']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 0)

categorical = ['workclass', 'education', 'marital.status',
'occupation', 'relationship', 'race', 'sex']
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])

X_train = pd.DataFrame(scaler.fit_transform(X_train), columns =
X.columns)

X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)

LR2 = LogisticRegression()
LR2.fit(X_train, y_train)

LogisticRegression()

y_pred = LR2.predict(X_test)
accuracy_score(y_test, y_pred)

0.8227044733340158

X = df.drop(['income', 'native.country', 'hours.per.week',
'capital.loss'], axis=1)
y = df['income']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 0)

categorical = ['workclass', 'education', 'marital.status',
'occupation', 'relationship', 'race', 'sex']
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])

X_train = pd.DataFrame(scaler.fit_transform(X_train), columns =
X.columns)

X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)

```

```

LR3 = LogisticRegression()
LR3.fit(X_train, y_train)

LogisticRegression()

y_pred = LR3.predict(X_test)

accuracy_score(y_test, y_pred)

0.8186098884225612

X = df.drop(['income'], axis=1)
y = df['income']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 0)

categorical = ['workclass', 'education', 'marital.status',
'occupation', 'relationship', 'race', 'sex', 'native.country']
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])

X_train = pd.DataFrame(scaler.fit_transform(X_train), columns =
X.columns)

pca= PCA()
pca.fit(X_train)
cumsum = np.cumsum(pca.explained_variance_ratio_)
dim = np.argmax(cumsum >= 0.90) + 1
print('The number of dimensions required to preserve 90% of variance
is',dim)

The number of dimensions required to preserve 90% of variance is 12

X = df.drop(['income', 'native.country', 'hours.per.week'], axis=1)
y = df['income']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.3, random_state = 0)
categorical = ['workclass', 'education', 'marital.status',
'occupation', 'relationship', 'race', 'sex']
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns =
X.columns)

```

```
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
LR2 = LogisticRegression()
LR2.fit(X_train, y_train)
```

```
LogisticRegression()
```

```
y_pred = LR2.predict(X_test)
accuracy_score(y_test, y_pred)
```

```
0.8227044733340158
```

```
from sklearn.metrics import confusion_matrix
import pandas as pd
confusion = confusion_matrix(y_test, y_pred)
df_confusion = pd.DataFrame(confusion, columns=['Predicted No',
'Predicted Yes'], index=['Actual No', 'Actual Yes'])
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
<=50K	0.84	0.95	0.89	7410
>50K	0.72	0.43	0.54	2359
accuracy			0.82	9769
macro avg	0.78	0.69	0.72	9769
weighted avg	0.81	0.82	0.81	9769



Conclusion:

Dimensionality reduction, specifically PCA, was applied to the Adult Census Income Dataset, aiming to simplify the dataset while preserving crucial information.

1. Original Dataset: Achieved 82.17% accuracy, and other metrics weren't provided.
2. Removing 'native.country': Slightly lower accuracy (82.13%) with comparable performance.
3. Removing 'hours.per.week': Slightly higher accuracy (82.27%) with similar performance.
4. Removing 'capital.loss': Slightly lower accuracy (81.86%) with similar performance.
5. PCA to retain 90% variance (12 dimensions): The dataset was reduced effectively while preserving essential information.

In this analysis, dimensionality reduction using PCA had a minor impact on model performance:

1. Accuracy: Slightly improved from 82.17% to 82.27%.
2. Precision: Improved slightly for the ">50K" class.
3. Recall: Decreased slightly for the ">50K" class.
4. F1 Score: Decreased slightly for the ">50K" class.